



LLM Validation: GPT-4o under High-Sycophancy Deployment

(A) Spiral Rates with 95% Bootstrap CIs



(B) Mean Final Beliefs vs LPC Threshold



Data: $n=200$ simulations per condition, $T=30$ turns, seeds=5000/5500/6000. High sycophancy prompt with GPT-4o.